

Project Scenario

In this assignment, you are a Data Analyst working at a Real Estate Investment Trust. The Trust would like to start investing in Residential real estate. You are tasked with determining the market price of a house given a set of features. You will analyze and predict housing prices using attributes or features such as square footage, number of bedrooms, number of floors, and so on. A template notebook is provided in the lab; your job is to complete the ten questions. Some hints to the questions are given in the template notebook.

Dataset Used in this Assignment

The dataset contains house sale prices for King County, which includes Seattle. It includes homes sold between May 2014 and May 2015. It was taken from [here](#). It was also slightly modified for the purposes of this course.

For this project, you will utilize JupyterLab running on the Cloud in Skills Network Labs environment.

Notebook URL: Alternatively, you can work on your local machine or any other environment of choice, by downloading this link : [Notebook link House Sales](#)



Launcher

House_Sales_in_King_Count_

Python (Pyodide)

```
[11]: df.dtypes

[11]: Unnamed: 0      int64
      id            int64
      date          object
      price         float64
      bedrooms      float64
      bathrooms     float64
      sqft_living    int64
      sqft_lot       int64
      floors        float64
      waterfront    int64
      view          int64
      condition     int64
      grade         int64
      sqft_above     int64
      sqft_basement  int64
      yr_built       int64
      yr_renovated   int64
      zipcode       int64
      lat           float64
      long          float64
      sqft_living15  int64
      sqft_lot15    int64
      dtype: object
```

We use the method describe to obtain a statistical summary of the dataframe.

↑

↓

↺

↻

🗑

Launcher X House_Sales_in_King_Count_ X + Python (Pyodide)

df.drop('id',axis=1,inplace=True)
df.drop('Unnamed: 0',axis=1,inplace=True)

df.describe()

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	co
count	2.161300e+04	21600.000000	21603.000000	21613.000000	2.161300e+04	21613.000000	21613.000000	21613.000000	21613.000000
mean	5.400881e+05	3.372870	2.115736	2079.899736	1.510697e+04	1.494309	0.007542	0.234303	0.234303
std	3.671272e+05	0.926657	0.768996	918.440897	4.142051e+04	0.539989	0.086517	0.766318	0.766318
min	7.500000e+04	1.000000	0.500000	290.000000	5.200000e+02	1.000000	0.000000	0.000000	0.000000
25%	3.219500e+05	3.000000	1.750000	1427.000000	5.040000e+03	1.000000	0.000000	0.000000	0.000000
50%	4.500000e+05	3.000000	2.250000	1910.000000	7.618000e+03	1.500000	0.000000	0.000000	0.000000
75%	6.450000e+05	4.000000	2.500000	2550.000000	1.068800e+04	2.000000	0.000000	0.000000	0.000000
max	7.700000e+06	33.000000	8.000000	13540.000000	1.651359e+06	3.500000	1.000000	4.000000	4.000000

We can see we have missing values for the columns bedrooms and bathrooms

```
[ ]: print("number of NaN values for the column bedrooms :", df['bedrooms'].isnull().sum())
print("number of NaN values for the column bathrooms :", df['bathrooms'].isnull().sum())
```

Launcher X House_Sales_in_King_Count_ Python (Pyodide)

Question 3


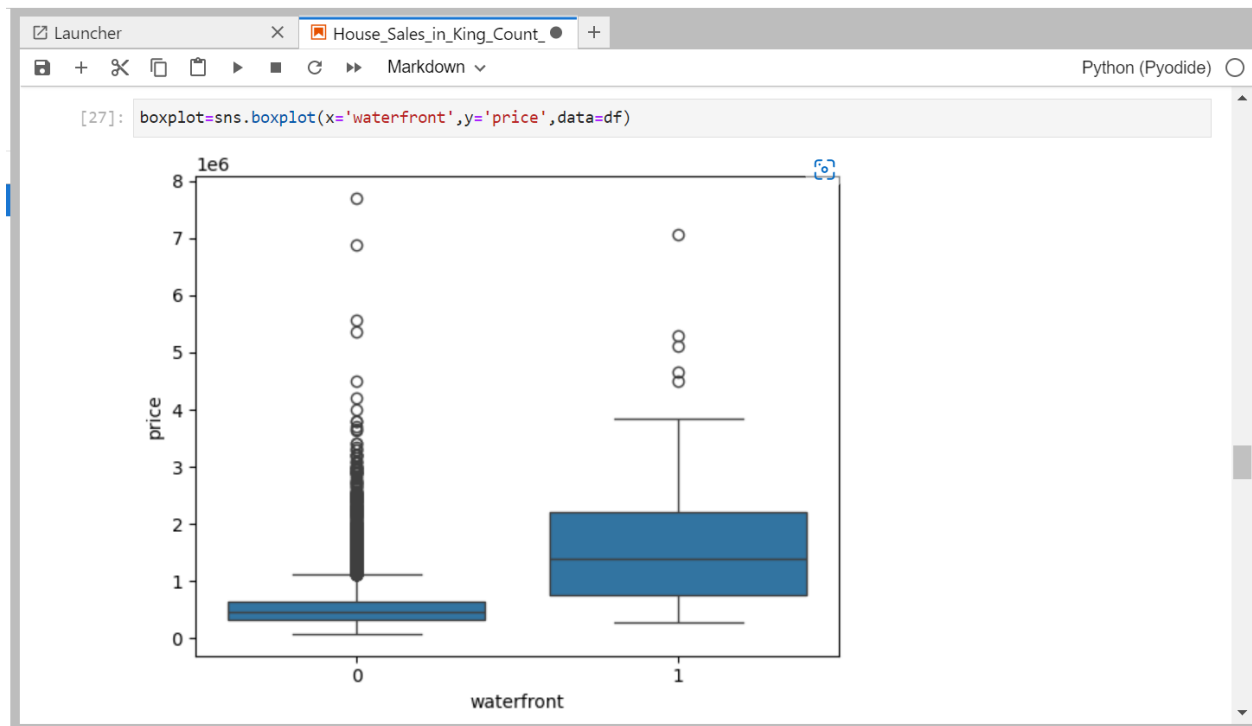
Use the method `value_counts` to count the number of houses with unique floor values, use the method `.to_frame()` to convert it to a dataframe.

```
[26]: df['floors'].value_counts().to_frame()
```

	floors
1.0	10680
2.0	8241
1.5	1910
3.0	613
2.5	161
3.5	8

Question 4

Use the function `boxplot` in the seaborn library to determine whether houses with a waterfront view or without a waterfront view have more price outliers.



Launcher x House_Sales_in_King_Count_ +

Python (Pyodide)

We can use the Pandas method `corr()` to find the feature other than price that is most correlated with price.

```
[29]: df.corr()['price'].sort_values()
```

```
[29]: zipcode      -0.053203
long           0.021626
condition      0.036362
yr_built       0.054012
sqft_lot15     0.082447
sqft_lot       0.089661
yr_renovated   0.126434
floors         0.256794
waterfront     0.266369
lat            0.307003
bedrooms       0.308797
sqft_basement  0.323816
view           0.397293
bathrooms      0.525738
sqft_living15  0.585379
sqft_above     0.605567
grade          0.667434
sqft_living    0.702035
price          1.000000
Name: price, dtype: float64
```

Module 4: Model Development

Launcher x House_Sales_in_King_Count_ +

Python (Pyodide)

Question 6

Fit a linear regression model to predict the 'price' using the feature 'sqft_living' then calculate the R^2 . Take a screenshot of your code and the value of the R^2 .

```
[31]: x1=df[['sqft_living']]
y1=df['price']
lr1=LinearRegression()
lr1.fit(x1,y1)
lr1.score(x1,y1)
```

```
[31]: 0.4928532179037931
```

Question 7

Fit a linear regression model to predict the 'price' using the list of features:

```
[ ]: features =["floors", "waterfront","lat" ,"bedrooms" ,"sqft_basement" ,"view" ,"bathrooms","sqft_living15","sqft_abov
```

Then calculate the R^2 . Take a screenshot of your code.

```
[ ]:
```


Launcher × House_Sales_in_King_Count_ × + Python (Pyodide)

```
print("number of test samples:", x_test.shape[0])
print("number of training samples:", x_train.shape[0])
```

number of test samples: 3242
number of training samples: 18371

Question 9

Create and fit a Ridge regression object using the training data, set the regularization parameter to 0.1, and calculate the R^2 using the test data.

```
[32]: from sklearn.linear_model import Ridge
```

```
[35]: rr=Ridge(alpha=0.1)
      rr.fit(x_train,y_train)
      rr.score(x_train,y_train)
```

```
[35]: 0.6594378534950235
```

Question 10

Perform a second order polynomial transform on both the training data and testing data. Create and fit a Ridge regression object using the training data, set the regularisation parameter to 0.1, and calculate the R^2 utilising the test data provided. Take a screenshot of your code and the R^2 .

Launcher × House_Sales_in_King_Count_ × + Python (Pyodide)

```
[35]: 0.6594378534950235
```

Question 10

Perform a second order polynomial transform on both the training data and testing data. Create and fit a Ridge regression object using the training data, set the regularisation parameter to 0.1, and calculate the R^2 utilising the test data provided. Take a screenshot of your code and the R^2 .

```
[36]: pt=PolynomialFeatures(degree=2)
      x_train_pt=pt.fit_transform(x_train)
      x_test_pt=pt.fit_transform(x_test)
      rr1=Ridge(alpha=0.1)
      rr1.fit(x_train_pt,y_train)
      rr1.score(x_test_pt,y_test)
```

```
[36]: 0.7002744263350642
```